

ImageNet CNN — SqueezeNet 1.1 spine, ImageNet (8 categories)

$n = 400$ stimuli, selected $K_{\max} = [5, 3, 2, 2, 3, 2, 5, 5, 10]$, threshold 0, 234-d fingerprint

selection: rank $\times 1.3$ of 6 swept configurations, silhouette 0.411 [0.399, 0.424] against 0.362 at the registry defaults, which is what every panel below is traced at
error bar in (g): paired 80 % stimulus subsample — this model has a single checkpoint, so no model-seed spread exists


(a) faithful: reconstruction controls




(b) reconstruction per node




(c) projection round-trip




(d) robust: stability across NMF seeds




(e) sensitivity to the rank




(f) rank vs. arbor explained



(g) vs. baselines: fingerprint or activations



(h) weights vs. activations only



(i) vs. pixel attribution

